

Build the Data Your AI Needs

Turn scattered information into trustworthy, reusable fuel for AI

A capable open model can explain semiconductor manufacturing, draft an invoice or imitate a customer-service conversation. Ask it about the warranty policy your company amended yesterday, however, and its confidence may outstrip its knowledge. The missing ingredient is not necessarily a bigger model. It is access to the right information, prepared in the right form.

This is the unglamorous engineering behind useful AI. Reports arrive as PDFs; rules hide in spreadsheets; useful answers are buried in emails; websites mix facts with advertising and navigation. Feed everything into an application indiscriminately and you have not created a knowledge base. You have created an expensive mess.

Our aim is to make a small, controlled collection of documents and examples that the projects ahead can reuse. We will prepare two distinct assets: a searchable knowledge corpus for retrieval-augmented generation (RAG), and an example dataset for teaching a model a particular